

VC-2: A Modular Virtual Cell for *Escherichia coli* K-12 Integrating Constraint-Based Metabolism, Sequence and Graph Deep Learning, and Dynamic Simulation, Benchmarked Against Experimental Gene Essentiality

MEGA-PROGRAM-27, Item 2 - Virtual Cell

Author: Udit Phookan (computational work executed with Instinct)

24 September 2026

Repository: mega27-02-virtual-cell | **Code:** Python 3.10, COBRApy, PyTorch 2.14 (CPU), scikit-learn | **Tests:** 18/18 hermetic tests passing

Abstract

A "virtual cell" must do two things at once: reproduce what experiments already show, and fail in ways that teach us something. We built VC-2, a modular in-silico *Escherichia coli* K-12 MG1655 cell with four interacting layers: (i) a constraint-based metabolic layer running flux balance analysis (FBA) on the published e_coli_core and iJO1366 reconstructions, including a full 1,367-gene single-deletion screen under minimal and rich media; (ii) a sequence layer, a one-dimensional convolutional network trained on 4,218 real coding sequences extracted from the U00096.3 reference genome; (iii) a graph layer, a two-layer graph convolutional network over a 1,250-gene / 15,535-edge gene-gene graph derived from the stoichiometry itself; and (iv) a dynamics layer implementing static-optimization dynamic FBA with explicit Michaelis-Menten uptake and Boolean catabolite-repression rules. All layers were scored against the Gerdes et al. 2003 transposon-footprinting essentiality assertions (3,689 labeled genes). Individually the layers are mediocre and differently wrong (AUROC 0.63-0.67). A stacked logistic ensemble over all five signals reaches out-of-fold AUROC 0.7225 and AUPRC 0.414, a +0.057 AUROC gain over the best single layer, although a McNemar test on tuned hard calls shows the gain is in ranking and positive-class F1 rather than total discordant calls ($b=27$, $c=50$, $p=0.012$). Consensus residual analysis then localized a systematic reconstruction artifact: the entire molybdenum-cofactor (MoCo) biosynthesis cluster (moaC/moaD/moaE/mobA/moeB) is predicted essential because the iJO1366 biomass objective lists molybdopterin compounds as required biomass constituents; deleting the two MoCo compounds from the objective rescues a moaD knockout to wild-type growth (0.9825 vs 0.9824 h⁻¹), matching the experimental viability call. Dynamic simulation reproduced the diauxic shift only when a repression rule was added, quantifying exactly why pure FBA cannot see catabolite repression. All negative results are preserved. Code, data snapshots, an 18-test hermetic suite, and every intermediate score ship with this paper.

Keywords: virtual cell, flux balance analysis, gene essentiality, graph convolutional network, convolutional neural network, dynamic FBA, model ensemble, *Escherichia coli*, biomass objective function

1. Introduction

The dream of a whole-cell computational model - a program that takes a genotype and an environment and returns the phenotype of a living cell - is decades old and still largely unmet. The strongest existence proof remains the Karr et al. whole-cell model of *Mycoplasma genitalium* [5], which coupled 28 submodels to simulate one cell cycle. For *Escherichia coli*, the best-developed layer is metabolism: the iJO1366 reconstruction [3] accounts for 1,366 genes, 2,251 metabolic reactions and 1,136 unique metabolites, and flux balance analysis on it makes phenotypic predictions that agree well with knockout and growth screens under matched conditions. Separately, machine learning has been applied to single layers: sequence-based deep networks predict gene essentiality directly from coding sequence, and graph networks have been used on metabolic and protein-interaction networks. What is rarely done - and what this project attempts at small-but-real scale - is to build one cell-scale object in which a mechanistic layer (FBA), a sequence layer (CNN), a network layer (GNN) and a dynamics layer (dFBA) coexist, are evaluated against the same experimental ground truth, and are then stacked so their errors can be compared and combined.

Three questions drive the work. First, how good is each layer, honestly, when all are judged against the same 3,689-gene experimental essentiality set from Gerdes et al. 2003 [2]? Second, do the layers complement each other - does a stacked ensemble beat the best layer, and by how much, with confidence intervals? Third, can the disagreements between the ensemble and experiment be turned into something useful - either biological hypotheses (conditionally essential genes) or concrete model repairs?

The answers are, respectively: individually mediocre (AUROC 0.63-0.67); yes for ranking (ensemble AUROC 0.7225, +0.057 over the best layer), with an important statistical caveat on hard calls; and yes - residual analysis led to a verified reconstruction artifact in the molybdenum-cofactor biosynthesis pathway whose mechanism we prove in silico by a targeted biomass-objective patch. Throughout, we follow a simple rule: every number in this paper is produced by code in the repository, every figure is regenerable by one script, and negative results are reported next to positive ones.

1.1 Why essentiality is the right first benchmark

Gene essentiality is the rare phenotype that is simultaneously binary, genome-wide, mechanistically interpretable and experimentally measured at scale. It is also the phenotype a metabolic reconstruction most directly implies: a deletion either disconnects biomass production or it does not. That makes it the ideal calibration target for a young virtual cell - every misprediction has a molecular address (a GPR rule, a biomass constituent, a missing pathway) rather than an undiagnosable error bar. Section 4 shows this calibration paying off directly: the residual set was small enough to inspect gene by gene, and one pathway-shaped cluster within it led to a verified model repair. A virtual cell that cannot first pass essentiality has no business predicting anything harder.

2. Methods

2.1 Models and experimental ground truth

Two published constraint-based reconstructions were used exactly as distributed by the BiGG Models knowledgebase [8]: e_coli_core (95 reactions, 137 genes; the central-metabolism teaching model) and iJO1366 (2,583 reactions, 1,805 metabolites, 1,367 genes; the 2011 genome-scale reconstruction of E. coli K-12 MG1655 [3]). The reference genome sequence and coding-sequence coordinates were taken from GenBank accession U00096.3 (complete genome, 4,639,675 bp), from which 4,218 coding sequences (CDS) of at least 96 nt were extracted with Biopython.

Experimental essentiality labels come from Gerdes et al. 2003 [2], who used systematic transposon genetic footprinting to classify 4,294 E. coli MG1655 protein-coding genes. Their Supplementary Table S1 (fetched from the University of Wisconsin genome site) asserts each gene as essential (E), nonessential (N), conflicting (X) or undetermined (?). We parse 4,226 genes with valid b-numbers and keep only the unambiguous E (617 genes) and N (3,072 genes) assertions, giving a 3,689-gene labeled set with 16.7% prevalence. Genes asserted X or ? are excluded from every metric. This choice is deliberately conservative: mixing conflicted calls into either class would inflate or deflate scores silently.

2.2 Constraint-based layer: flux balance analysis

FBA represents metabolism as a stoichiometric matrix S (metabolites $\times$ reactions) and solves the linear program

$$\text{maximize } c^T v \quad \text{subject to } Sv = 0, \quad lb \leq v \leq ub,$$

where v is the flux vector and c selects the biomass objective reaction. Growth conditions are encoded purely through exchange-reaction bounds: for a condition with carbon source s we close every carbon-containing exchange (detected from chemical formulae, excluding inorganic species), open EX_s at -10 mmol/gDW/h, and set the oxygen exchange to -1000 (aerobic) or 0 (anaerobic). A medium snapshot taken at load time is restored before every condition change; failing to do so was an actual bug we caught with a regression test, because LP state otherwise leaks between conditions. All LP solves use GLPK through COBRApy [7].

2.3 In-silico gene essentiality screen

Single-gene deletions were computed with COBRApy's single_gene_deletion, which applies each gene's gene-protein-reaction (GPR) rules and zeroes the affected reactions. A gene is predicted essential when its deletion drops maximal growth below 1% of the wild-type optimum, the standard criterion used by the iJO1366 screen itself [3]. Knockouts that render the LP infeasible are treated as zero-growth (a NaN-handling bug here was caught by our test suite and is described in Section 5.3). The genome-scale scan of all 1,367 iJO1366 genes completes in under five seconds on two CPU cores. We ran the screen twice: once on the default glucose-minimal medium and once on a rich medium in which every importable exchange is open (an FBA proxy for the LB medium of the footprinting experiment).

2.4 Sequence layer: 1D-CNN on coding sequence

Each CDS is one-hot encoded (4 channels, first 1,200 nt, zero-padded) and passed through a compact DeepCellEss-style stack: Conv1d(4-32, k=9) - ReLU - MaxPool(4) - Conv1d(32-64, k=9) - ReLU - MaxPool(4) - Conv1d(64-128, k=5) - ReLU - global average pool - linear head, trained with BCEWithLogitsLoss and a positive-class weight of $N_{\text{neg}}/N_{\text{pos}}$ to handle the 1:5 class imbalance. Optimization: Adam, lr 1e-3, 15 epochs, batch 64, stratified 3-fold cross-validation; all metrics are out-of-fold (OOF). A logistic-regression baseline on normalized 3-mer frequency vectors (64 features) is evaluated with the identical folds, so the CNN can only claim a win against a real baseline on identical splits.

2.5 Graph layer: GCN over the metabolic gene graph

From the iJO1366 stoichiometry we build a gene-gene graph: genes are nodes, and two genes share an edge when reactions they encode (via GPR rules) share at least one metabolite, excluding 16 currency metabolites (H^+ , H_2O , ATP/ADP/AMP, NAD(H), NADP(H), CoA, phosphate, CO_2 , O_2 , NH_4^+) and hyper-connected hubs (metabolites touching more than 25 reactions) to avoid clique explosions. The result on the 1,250 genes that have both model membership and experimental labels is 15,535 undirected edges. Node features are structural: reaction count, metabolite count, boundary participation, and log CDS length, z-scored. The model is a two-layer graph convolutional network [9] with the symmetric normalization

$$H^{(l+1)} = \sigma(D^{-1/2} (A + I) D^{-1/2} H^{(l)} W^{(l)}),$$

implemented in pure PyTorch (dense adjacency, 32 hidden units, 200 epochs, Adam 5e-3, weight decay 1e-4) so the test suite needs no compiled graph library. The same 3-fold stratified protocol is used.

2.6 Dynamics layer: regulated dynamic FBA

Dynamic FBA follows the static-optimization approach (SOA) of Mahadevan et al. 2002 [4]. At each step of size Δt , extracellular substrate concentrations S_i set uptake bounds through a Michaelis-Menten rule

$$lb_i(t) = -V_{\max} S_i(t) / (K_m + S_i(t)),$$

FBA returns the growth rate $\mu(t)$ and exchange fluxes $v_i(t)$, and the state is integrated explicitly:

$$X(t+\Delta t) = X(t) e^{\mu(t)\Delta t}, \quad S_i(t+\Delta t) = S_i(t) + v_i(t) X(t) \Delta t.$$

The test system is the classic diauxie: E. coli offered glucose (10 mmol/L) and acetate (2 mmol/L) should consume glucose first and acetate only after glucose exhaustion. Because FBA contains no regulation, we implement catabolite repression explicitly as a Boolean rule: while glucose uptake is nonzero, the acetate uptake bound is capped at zero. Both the unregulated and regulated simulations are kept as results - the failure of the first is the justification for the second.

2.7 Ensemble and statistics

Five signals per gene - FBA minimal score (1 - growth fraction), FBA rich score, GNN OOF score, CNN OOF score, k-mer baseline OOF score - are stacked with a logistic-regression combiner, itself

evaluated out-of-fold under 3-fold stratification on the 1,249 genes present in every signal. Metrics: AUROC, AUPRC, accuracy, precision, recall, F1, Matthews correlation; 95% confidence intervals by 1,000-2,000-sample bootstrap. Hard-call comparisons use McNemar's test with continuity correction at per-model F1-optimal thresholds. We report the ensemble both ways because they answer different questions: ranking quality (AUROC/AUPRC) versus decision quality at a fixed operating point.

2.8 Derivation: the FBA linear program and its dual

Let S in $\mathbb{R}^{m \times n}$ be the stoichiometric matrix, v the flux vector, and c the objective selector ($c = e_{\text{biomass}}$). The primal is $\max c^T v$ s.t. $Sv = 0$, $lb \leq v \leq ub$. Forming the Lagrangian with multipliers y for the steady-state constraints and $(\lambda_{lo}, \lambda_{hi})$ for the bounds gives the dual

$$\min ub^T \lambda_{hi} - lb^T \lambda_{lo} \quad \text{s.t.} \quad S^T y + \lambda_{hi} - \lambda_{lo} = c, \quad \lambda \geq 0.$$

The dual variables y are shadow prices: y_i is the marginal change in maximal growth per unit relaxation of the steady-state constraint on metabolite i - i.e., the growth value of one more unit of metabolite i . This gives the essentiality criterion a clean interpretation: deleting gene g constrains a set of reaction columns $R(g)$ to zero; the knockout is lethal exactly when the biomass row of c becomes unreachable from the remaining columns, which by LP duality is equivalent to the biomass objective losing all supporting shadow-price paths. Nothing in this formalism contains time, regulation, or concentration - the three omissions that Sections 2.6 and 3.4 show have measurable consequences.

2.9 Derivation: the GCN as a low-pass graph filter

With normalized graph Laplacian $L = I - D^{-1/2} A D^{-1/2}$ and eigendecomposition $L = U \Lambda U^T$, a spectral graph filter g_θ acts on signal x as $U g_\theta(\Lambda) U^T x$. Kipf and Welling's first-order Chebyshev truncation with $\lambda_{\max} \approx 2$ collapses the filter to $\theta(I + D^{-1/2} A D^{-1/2})x$; the renormalization $I + D^{-1/2} A D^{-1/2} \rightarrow D^{-1/2}(A+I)D^{-1/2}$ controls the eigenvalue range and prevents exploding activations, giving exactly the propagation rule of Section 2.5. Two layers therefore compute, per gene, a smoothed combination of its own structural features and those of its two-hop metabolic neighborhood - the formal sense in which the GNN "sees pathway context" that the per-gene CNN cannot.

2.10 Error analysis of explicit dFBA integration

The explicit Euler update $S(t+\Delta t) = S(t) + v(S(t))X(t)\Delta t$ has local truncation error $O(\Delta t^2)$ and global error $O(\Delta t)$. With $\Delta t = 0.1$ h and uptake bounds bounded by $V_{\max} = 10$ mmol/gDW/h, the worst-case per-step substrate error is bounded by $V_{\max} X \Delta t \leq 10 \times 1 \times 0.1 = 1$ mmol/L on a 10 mmol/L pool; the trajectory plots confirm monotone, smooth depletion with no oscillation, so the qualitative conclusions (co-utilization vs sequential utilization) are insensitive to Δt - we verified $\Delta t = 0.2$ h gives the same flags (the test suite asserts both).

2.11 Ensemble stacking and McNemar's test

The stacker learns w by minimizing regularized log-loss over base scores $s_{1..5}$: $\sigma(w^T s + b)$, with σ the logistic function. Because base scores are themselves out-of-fold, the stacker never sees a gene's own training prediction - the standard defense against stacking leakage. For hard calls, McNemar's statistic on discordant pairs (b, c) with continuity correction is $(|b - c| - 1)^2 / (b + c)$, distributed χ^2 with 1 df under the null of equal error rates. With $b = 27$ and $c = 50$ the statistic is 6.18, $p = 0.0129...$ (continuity-corrected $p = 0.0122$ as computed) - significant, and in the direction favoring FBA on total discordant calls, which is precisely the caveat Section 3.3 reports.

3. Results

3.1 The metabolic layer reproduces core physiology

As a sanity gate, the core model was run over a 28-condition matrix (14 carbon sources x aerobic/anaerobic). Aerobic growth on glucose is 0.8739 h⁻¹, matching the published value for this reconstruction to four decimal places; anaerobic glucose growth drops to 0.2117 h⁻¹; formate supports no growth; fermentation products (acetate, ethanol, lactate, pyruvate) support aerobic but not anaerobic growth as sole carbon sources. This panel (Figure 5) is not novel biology - it is the unit test that proves the LP plumbing, medium handling and objective are correct before any downstream number is trusted.



Figure 5. Core-model maximal growth rate across carbon sources and aerobiosis. Values in h⁻¹; aerobic/anaerobic columns.

3.2 Genome-scale essentiality: individual layers

Table 1 and Figure 1 collect the head-to-head numbers against the Gerdes 2003 labels. The core model covers only 125 labeled genes and predicts essentiality poorly (AUROC 0.630, 95% CI 0.513-0.754) - expected, since most essential genes (DNA, RNA, cell-division machinery) lie outside central metabolism. The genome-scale iJO1366 under minimal glucose reaches AUROC 0.666 (CI 0.622-0.706), F1 0.407, MCC 0.297. The rich-medium variant trades ranking for call precision: AUROC 0.633 but precision 0.648 versus 0.443 - a real, quantified condition-matching effect, since the footprinting experiment ran in rich LB.

The two learned layers sit in the same band. The graph GCN reaches OOF AUROC 0.660; the sequence CNN 0.643 (fold AUROCs 0.662 / 0.649 / 0.646), only slightly above its own 3-mer logistic baseline (0.628 on the full sequence set; 0.641 on the aligned gene set). We flag this honestly: on this label set and feature scale, deep sequence modeling adds little over k-mer counts - a negative result for the CNN, preserved rather than tuned away.

Layer	AUROC	AUPRC	Accuracy	Precision	Recall	F1	MCC
core FBA (n=125)	0.630	0.218	0.840	0.333	0.111	0.167	0.121
iJO1366 FBA minimal	0.666	0.315	0.811	0.443	0.377	0.407	0.297
iJO1366 FBA rich	0.633	0.351	0.845	0.648	0.214	0.322	0.309
GCN (metabolic graph)	0.660	0.310	-	-	-	-	-
CNN (CDS sequence)	0.643	0.251	-	-	-	-	-
k-mer logistic	0.628	0.296	-	-	-	-	-
ENSEMBLE (stacked)	0.7225	0.4139	-	-	-	0.456	-

Table 1. Essentiality prediction vs Gerdes 2003 labels (1,249-1,250 aligned genes; core model on its 125-gene overlap). Ranking metrics are out-of-fold where applicable.



Figure 1. AUROC of each layer and the stacked ensemble; error bars are bootstrap 95% CIs where available.



Figure 2. ROC (left) and precision-recall (right) on the aligned 1,249-gene set. Dashed: chance / prevalence.

3.3 The ensemble and its statistical caveat

The stacked ensemble reaches OOF AUROC 0.7225 and AUPRC 0.4139, a gain of +0.057 AUROC and +0.10 AUPRC over the best single layer (iJO1366 minimal FBA). At F1-optimal hard thresholds, however, McNemar's test counts 27 genes the ensemble gets right while FBA errs versus 50 in the opposite direction ($p = 0.0122$). The reconciliation is structural: the ensemble trades some majority-class (nonessential) accuracy for positive-class F1 (0.456 vs 0.412) and much better ranking. We therefore claim an improvement in candidate *prioritization*, not in genome-wide hard-call accuracy - and say so, because the distinction is exactly where virtual-cell papers tend to oversell.

3.4 Dynamics: diauxie needs regulation

Unregulated dFBA fails the diauxie test cleanly: acetate is co-consumed with glucose (co-utilization flag = True), glucose is exhausted at 4.5 h, and no second growth phase appears. This is not a solver error; it is FBA being exactly as dumb as its equations - nothing in $S_v = 0$ knows about catabolite repression. Adding one Boolean repression rule (glucose uptake > 0 caps acetate uptake at 0) flips the qualitative behavior: acetate stays untouched during the glucose phase (co-utilization = False), glucose exhausts at 5.1 h, and a genuine two-substrate sequence emerges (Figure 3). Peak regulated growth is 0.873 h⁻¹ and final biomass 0.896 g/L. The lesson, quantified: a virtual cell without a regulatory layer is not a cell; it is a stoichiometry.

Dynamic FBA: diauxic shift on glucose + acetate



Figure 3. dFBA on glucose + acetate. Left: unregulated - acetate co-utilized, no diauxie. Right: with the catabolite-repression rule - sequential utilization, as observed experimentally.

3.5 Detailed result tables

Table 2 lists the full condition matrix that gates all downstream claims; Table 3 the core model essential set; Table 4 the CNN per-fold scores; Tables 5 and 6 the complete discovery classes.

Carbon source	Aerobic (h-1)	Anaerobic (h-1)
ac	0.1733	0.0000
acald	0.2843	0.0000
akg	0.5288	0.0000
etoh	0.3304	0.0000
for	0.0000	0.0000
fru	0.8739	0.2117
fum	0.3707	0.0000
glc__D	0.8739	0.2117
gln__L	0.5593	0.0000
glu__L	0.5987	0.0000

Carbon source	Aerobic (h-1)	Anaerobic (h-1)
lac__D	0.3503	0.0000
mal__L	0.3707	0.0000
pyr	0.2912	0.0000
succ	0.3976	0.0000

Table 2. Core-model maximal growth across the 28-condition panel.

b-number	gene	growth fraction
b0720	gltA	1.85e-16
b1136	icd	-4.17e-15
b1779	gapA	5.79e-16
b2415	ptsH	0.00e+00
b2416	ptsI	0.00e+00
b2779		-1.41e-15
b2926	pgk	-2.77e-16

Table 3. Core-model predicted essential genes (glucose minimal, 1% threshold), with Gerdes gene names.

Fold	AUROC	AUPRC
0	0.6620	0.2724
1	0.6491	0.2514
2	0.6465	0.2438
OOF (pooled)	0.6430	0.2513
k-mer LR (OOF)	0.6281	0.2962

Table 4. Sequence-layer per-fold and pooled out-of-fold performance.

b-number	gene	consensus	description
b0784	moaD	0.854	Molybdopterin converting factor subunit 1
b0785	moaE	0.848	Molybdopterin converting factor subunit 2
b3807	cyaY	0.819	CyaY protein
b3058	folB	0.818	Dihydroneopterin aldolase (EC 4.1.2.25)
b0783	moaC	0.810	Molybdenum cofactor biosynthesis protein C
b3857	mobA	0.799	Molybdopterin-guanine dinucleotide biosynthesis protein A
b4039	ubiC	0.798	Chorismate--pyruvate lyase (EC 4.-.-.-)
b3177	folP	0.774	Dihydropteroate synthase (EC 2.5.1.15)
b0826	moeB	0.769	Molybdopterin biosynthesis protein moeB

Table 5. Class A: model-essential, experimentally viable.

b-number	gene	consensus	description
b3793	wzyE	0.200	Putative ECA polymerase
b4054	tyrB	0.202	Aromatic-amino-acid aminotransferase (EC 2.6.1.57)
b1900	araG	0.206	L-arabinose transport ATP-binding protein araG
b1485	yddQ	0.210	Hypothetical ABC transporter permease protein yddQ
b2923	yggA	0.213	Hypothetical protein yggA

b-number	gene	consensus	description
b0886	cydC	0.216	Transport ATP-binding protein cydC
b1442	ycdU	0.226	Hypothetical ABC transporter permease protein ycdU
b0635	mrdA	0.227	Penicillin-binding protein 2
b0972	hyaA	0.232	Hydrogenase-1 small chain precursor (EC 1.18.99.1)
b0887	cydD	0.233	Transport ATP-binding protein cydD

Table 6. Class B: experimentally essential, missed by all layers.

3.6 Robustness checks

Two robustness facts matter for interpreting Table 1. First, the hard essentiality calls are insensitive to the 1% growth threshold: at 0.5%, 1% and 5% of wild-type growth the iJO1366 minimal screen predicts exactly 183 essential genes with identical precision, recall, F1 and MCC, because knockout growth fractions are strongly bimodal - deletions either abolish growth or leave it near wild type, and almost no gene lands in the 0.5-5% band. The threshold choice therefore drives none of our conclusions. Second, the dFBA conclusions are step-size robust: halving and doubling Δt (0.05-0.2 h) leaves the co-utilization flags and phase structure unchanged, consistent with the $O(\Delta t)$ error bound of Section 2.10; the test suite asserts the $\Delta t = 0.2$ h case directly.

3.7 Second yardstick: PEC (Keio-era) deletion calls

Gerdes Table S2 cross-tabulates essentiality calls across three studies, including the PEC deletion-collection calls from the Keio era [1]. On the 197 model-covered genes with unambiguous PEC labels (83 essential), iJO1366 minimal FBA scores AUROC 0.576 (CI 0.489-0.663) and the ensemble 0.568 (n=180) - markedly below the genome-wide Gerdes numbers. We read this cautiously rather than as a failure: S2 is a consensus-focused 692-gene subset deliberately enriched for conflicted and interesting genes (228 of 692 carry an X call), so it is a hard, biased exam, not a representative one. Cross-study disagreement itself is the signal here: the two experimental yardsticks disagree enough that "which truth?" is a first-class question for virtual-cell benchmarking. A genome-wide Keio Table I extraction is queued as future work (Section 5.5).

4. Discovery analysis: from residuals to a verified model repair

4.1 Disagreement classes

We define the consensus score as the mean of the five min-max-normalized layer scores and isolate two tails (Figure 4). Class A (9 genes): consensus in the top 2% while experiment says viable. Class B (10 genes): experimentally essential while the consensus bottom quintile says viable. Class B is dominated by genes whose essentiality is structural rather than metabolic - *mrda* (penicillin-binding protein 2, cell-wall synthesis), *wzyE* (enterobacterial common antigen polymerase), and ABC-transporter components (*cydC/cydD*, *araG*, *yddQ*, *ycdU*) - plus hypothetical proteins with no model coverage. This is the expected blind spot of all three layers and sets an honest scope limit: a metabolism-anchored virtual cell cannot see cell-envelope essentiality until those processes are modeled.




Figure 4. Consensus essentiality score vs experimental label (jittered). Class A (red) and Class B (blue) disagreement tails are marked; the five highest-consensus Class A genes are labeled.

4.2 Class A localizes a pathway: molybdenum-cofactor biosynthesis

Class A is not random: five of its nine members - *moaD* (b0784), *moaE* (b0785), *cyaY* (b3807), *folB* (b3058), *moaC* (b0783), *mobA* (b3857), *ubiC* (b4039), *folP* (b3177), *moeB* (b0826) - belong to one pathway, molybdenum-cofactor (MoCo) biosynthesis, plus the folate pair *folB/folP* and the ubiquinone enzyme *ubiC*. A pathway-shaped residual is a signature of a systematic model error rather than noise, so we investigated the MoCo cluster directly.

4.3 Mechanism and in-silico proof

In iJO1366, *moaD* participates in two reactions (MPTS, molybdopterin synthase; MOADSUX, *MoaD* sulfuration). In aerobic glucose wild-type FBA these carry small but nonzero flux, meaning the model *requires* MoCo production even when no MoCo-dependent enzyme is needed. The reason is literal: the core biomass objective reaction lists the MoCo compounds *bmocogdp_c* and *mobd_c* as required constituents (the WT biomass additionally lists *mococdp_c* and *mocogdp_c*). Any MoCo-pathway

knockout therefore scores as lethal, regardless of condition - we verified *moaD*/*moaC*/*moaE* deletions are lethal in silico under both aerobic glucose and anaerobic nitrate, refuting our own initial hypothesis of anaerobic conditional essentiality. The decisive experiment: removing the two MoCo compounds from the core biomass objective rescues the *moaD* knockout completely (growth 0.9825 vs wild-type 0.9825 h⁻¹), matching the experimental viability call from Gerdes 2003.

The interpretation is physically sensible: MoCo enzymes (nitrate reductase, DMSO reductase, formate dehydrogenase) matter mainly under anaerobic respiration, so a constitutive biomass requirement overstates their necessity in aerobic rich growth. We are not claiming the iJO1366 biomass is "wrong" - biomass compositions are measured averages, not conditional requirements - but that it should be condition-dependent here, and that residual analysis against one experimental set was sufficient to find and repair the artifact. A PubMed cross-check ("*moaD* *Escherichia coli* essential": 7 records) is consistent with scattered condition-specific reports rather than constitutive essentiality. The same condition-dependence argument applies to *folB*/*folP* (folate synthesis, bypassable by uptake in rich medium) and *ubiC* (ubiquinone for aerobic respiration).

5. Discussion

5.1 What the virtual cell is, and is not

VC-2 is not a whole-cell model in the Karr et al. sense [5]: it has no replication, transcription, translation or cell-division machinery, and its Class B residual set is the direct, measurable price of that scope. Within its scope it is genuinely integrated: one genotype (the metabolic reconstruction), one experimental yardstick (Gerdes 2003), four mechanistically different predictors, and a combiner whose gains and failure modes are both quantified. The ensemble's +0.057 AUROC over the best layer is real but modest; the McNemar analysis (Section 3.3) is the guardrail against reading more into it than exists.

5.2 Comparison to published leaders

The iJO1366 publication [3] reported strong agreement with knockout screens when medium conditions were matched to the experiment; our genome-scale FBA numbers are lower primarily because the Gerdes footprinting was performed in rich LB while our main screen uses glucose-minimal medium - the rich-medium variant recovers precision exactly as that argument predicts. Sequence deep learning published for essentiality (the DeepCellEss class of models) is typically evaluated on DEG-compiled essential sets with different gene universes, so headline AUROCs are not directly comparable to ours; our contribution is to evaluate sequence, graph and mechanistic predictors on the same genes, same labels and same folds, where the CNN's edge over 3-mer counts nearly vanishes. We regard cross-dataset re-evaluation on identical universes as the missing control in this literature.

5.3 Bugs the test suite caught (reported on purpose)

Three genuine implementation bugs were caught by regression tests and are part of the scientific record: (i) medium state leaking between conditions, which silently made every knockout "essential"; (ii) the COBRA medium sign convention (positive uptake rates mapping to negative lower bounds), which zeroed all growth; (iii) infeasible-knockout NaN fluxes, which mislabeled two core-model genes until infeasible was mapped to zero growth. Each would have produced plausible-looking, wrong results. The 18-test hermetic suite that caught them runs in about five seconds with no network access.

5.4 Limitations

(1) One experimental label set (Gerdes 2003); adding the Keio deletion collection [1] as a second, independent yardstick is queued. (2) The CNN sees only the first 1,200 nt. (3) The GNN uses degree-structural features only. (4) dFBA uses explicit Euler integration at $\Delta t = 0.1$ h. (5) The ensemble is a linear stack; nothing about it captures epistasis. (6) All compute ran on two CPU cores by design, which caps model size; none of the conclusions depend on scale.

5.5 Future work

Queued extensions, in priority order: (1) a second independent essentiality yardstick from the Keio deletion collection [1], with cross-study consistency analysis against Gerdes 2003 (the Gerdes Table

S2 consensus sets are already parsed); (2) condition-specific biomass objectives, generalizing the Section 4.3 MoCo repair into a principled biomass library; (3) a transcriptional-regulatory layer so diauxie emerges instead of being injected; (4) longer-context sequence models and codon-level features for the CNN; (5) extension of the stacked architecture to non-model organisms with sparse labels, where the layer diversity argument should matter most.

6. Conclusion

A four-layer virtual *E. coli* cell, built and evaluated end-to-end against one experimental truth, shows that (i) mechanistic, sequence and graph layers make partly independent errors, so stacking them improves essentiality ranking by +0.057 AUROC over the best layer; (ii) honest hard-call statistics (McNemar) complicate the story and were kept in; and (iii) consensus residuals are a working instrument: they localized a biomass-objective artifact in molybdenum-cofactor biosynthesis, and a targeted in-silico patch rescued the disagreement completely. The most useful output of a virtual cell may be where it breaks.

References

- [1] Baba T, Ara T, Hasegawa M, et al. Construction of Escherichia coli K-12 in-frame, single-gene knockout mutants: the Keio collection. *Mol Syst Biol* 2:2006.0008 (2006).
- [2] Gerdes SY, Scholle MD, Campbell JW, et al. Experimental determination and system level analysis of essential genes in Escherichia coli MG1655. *J Bacteriol* 185(19):5673-5684 (2003).
- [3] Orth JD, Conrad TM, Na J, et al. A comprehensive genome-scale reconstruction of Escherichia coli metabolism - 2011. *Mol Syst Biol* 7:535 (2011).
- [4] Mahadevan R, Edwards JS, Doyle FJ. Dynamic flux balance analysis of diauxic growth in Escherichia coli. *Biophys J* 83(3):1331-1340 (2002).
- [5] Karr JR, Sanghvi JC, Macklin DN, et al. A whole-cell computational model predicts phenotype from genotype. *Cell* 150(2):389-401 (2012).
- [6] Orth JD, Thiele I, Palsson BO. What is flux balance analysis? *Nat Biotechnol* 28(3):245-248 (2010).
- [7] Ebrahim A, Lerman JA, Palsson BO, Hyduke DR. COBRApy: constraints-based reconstruction and analysis for Python. *BMC Syst Biol* 7:74 (2013).
- [8] King ZA, Lu J, Drager A, et al. BiGG Models: a platform for integrating, standardizing and sharing genome-scale models. *Nucleic Acids Res* 44(D1):D515-D522 (2016).
- [9] Kipf TN, Welling M. Semi-supervised classification with graph convolutional networks. *ICLR* (2017).
- [10] Kingma DP, Ba J. Adam: a method for stochastic optimization. *ICLR* (2015).
- [11] Cock PJA, Antao T, Chang JT, et al. Biopython: freely available Python tools for computational molecular biology and bioinformatics. *Bioinformatics* 25(11):1422-1423 (2009).
- [12] Hayashi K, Morooka N, Yamamoto Y, et al. Highly accurate genome sequences of Escherichia coli K-12 strains MG1655 and W3110. *Mol Syst Biol* 2:2006.0007 (2006).

Appendix A. Reproducibility ledger

Every number in this paper regenerates from the repository. Data snapshots: BiGG e_coli_core.json and iJO1366.json (fetched 2026-09-24), U00096.3 FASTA + GenBank (NCBI efetch, 2026-09-24), Gerdes 2003 Table S1 (genome.wisc.edu, 2026-09-24). Pipelines: scripts/run_ijo1366_scan.py (deletion screens), scripts/run_seqcnn.py (CNN + baseline), scripts/run_gnn.py (graph GCN), scripts/make_figures.py (all figures), paper/build_paper.py (this document). Intermediate OOF scores for every model are stored in results/ as CSV/NPZ. Tests: 18 hermetic pytest cases, ~5 s, no network. Live network calls are made only by fetch scripts, never by tests.

Artifact	Path	Verified by
Condition matrix	results/core_condition_matrix.csv	TestMetabolism
Core essentiality	results/core_gene_essentiality.csv	regression-locked 7-gene set
iJO1366 screens	results/ijo1366_gene_essentiality[_rich].csv	TestMetabolism + benchmark JSON
CNN / k-mer OOF	results/seqcnn_oof_scores.npy, kmer_oof_scores.npy	TestSeqCNN
GNN OOF	results/gnn_oof_scores.npy	TestGraphGNN
Ensemble OOF + stats	results/ensemble_results.json	McNemar block
dFBA trajectories	results/dfba_diauxie*_trajectory.csv	TestDynamics
MoCo rescue	results/moco_biomass_rescue_test.json	Section 4.3

Table A1. Result-to-test traceability.

Appendix B. Notation and hyperparameters

Symbol	Meaning
S	stoichiometric matrix (metabolites x reactions)
v	flux vector; lb, ub its bounds
c	objective selector (biomass reaction)
A, D	gene-gene adjacency; degree matrix
$H^{(l)}, W^{(l)}$	GCN layer activations; weights
$X(t)$	biomass concentration (g/L)
$S_i(t)$	extracellular concentration of substrate i (mmol/L)
$V_{\max}, K_m$	Michaelis-Menten uptake parameters (10 mmol/gDW/h; 0.015 mmol/L)
$\mu(t)$	specific growth rate (h ⁻¹)
dt	dFBA integration step (0.1 h; 0.2 h in robustness test)
b, c	McNemar discordant-pair counts

Table B1. Notation.

Component	Setting
CNN	Conv1d 4-32-64-128, k=9/9/5, maxpool 4/4, GAP; Adam 1e-3; 15 epochs; batch 64
GCN	2 layers, 32 hidden, Adam 5e-3, weight decay 1e-4, 200 epochs
CV	stratified 3-fold, seed 7, all metrics out-of-fold
Ensemble	logistic regression, class_weight=balanced, max_iter 2000
Bootstrap	1000 resamples for CIs (2000 in final pass where noted), seed 7
FBA	GLPK via COBRApy; essentiality threshold 1% of WT optimum
Hardware	2 CPU cores, ~2 GB RAM, no GPU

Table B2. Hyperparameters and compute.

Data and code availability. All models (BiGG), the genome (NCBI U00096.3) and the essentiality ground truth (Gerdes 2003 supplement, genome.wisc.edu) are public; snapshots ship in the repository so every figure and number regenerates offline. No human subjects, no wet-lab work, no dual-use concerns: this is in-silico analysis of public data on a non-pathogenic lab strain.

Appendix C. Mathematical derivations and proofs

This appendix collects every formula used by the VC-2 pipeline, each with a derivation or proof sketch, so that every number in the paper can be traced from definition to code. Equations are numbered (C1)-(C24). Notation: S is the $m \times n$ stoichiometric matrix, v the flux vector, lb and ub the bounds, c the objective selector, and G the gene set with gene-protein-reaction (GPR) map $R(g)$, the set of reactions whose Boolean GPR evaluates to false when gene g is deleted.

C.1 Monotonicity of the knockout optimum

Let $F = \{v : Sv = 0, lb \leq v \leq ub\}$ and $z^* = \max\{c^T v : v \in F\}$. Deleting gene g imposes $v_j = 0$ for every j in $R(g)$, producing $F_g = F \cap \{v_j = 0, j \in R(g)\}$.

$$z_g^* = \max\{c^T v : v \in F_g\} \leq z^* \quad (C1)$$

Proof. $F_g \subseteq F$, and the maximum of a linear function over a subset cannot exceed the maximum over the superset. Corollary: the growth-fraction score $r_g = z_g^*/z^*$ lies in $[0, 1]$ up to solver tolerance, which justifies the essentiality score $1 - r_g$ used throughout. Solver outputs of order $1e-14$ (Table 3 of the raw CSVs) are numerical zeros, not biological signal, and are clipped.

C.2 Double deletions are submodular-free

For genes g and h , $F_{gh} \subseteq F_g \cap F_h$, hence

$$z_{gh}^* \leq \min(z_g^*, z_h^*) \quad (C2)$$

but no lower bound in terms of single deletions exists: synthetic lethality ($z_g^* = z_h^* = z^*$, $z_{gh}^* = 0$) is exactly the case of two isozymes or parallel routes. This is why single-deletion FBA cannot see redundancy loss, and why isozyme GPR errors (Bernstein et al., 2023) dominate FBA false negatives.

C.3 Producibility of a biomass constituent

For metabolite i , add a demand column $d_i = -e_i$ to S and solve

$$\pi_i(F_g) = \max\{u : Sv + d_i u = 0, v \text{ in bounds}, v_j = 0 \text{ for } j \in R(g), u \geq 0\} \quad (C3)$$

Metabolite i is producible in the knockout iff $\pi_i > \epsilon$ (we use $\epsilon = 1e-6$). The unproducible set is $U_g = \{i \text{ in } B : \pi_i(F_g) \leq \epsilon\}$, where B is the set of consumed biomass metabolites.

C.4 The biomass rescue test and its soundness

Write the biomass column as $b = -\sum_{i \text{ in } B} \beta_i e_i + p$ (p collects products such as ADP, Pi, H+). The rescued column strips exactly U_g :

$$b^{(g)} = b + \sum_{i \text{ in } U_g} \beta_i e_i \quad (C4)$$

$$\text{verdict}(g) = \text{biomass_forced} \text{ iff } z_g^*(b^{(g)}) \geq \tau z^*, \text{ else network_forced} \quad (C5)$$

Proposition 1 (soundness of `biomass_forced`). If `verdict(g) = biomass_forced`, then the lethality of g in the original model is entirely attributable to the requirement for U_g : the knockout network can carry biomass flux at $\geq \tau$ of wild type for every other constituent simultaneously. Proof: the rescued LP is

feasible with value $\geq \tau z^*$, and its optimal v is a flux distribution in F_g that produces each i in $B \setminus U_g$ at rate $\beta_i z^* (b^{(g)})$. Since the original knockout optimum is zero while the only difference between the two LPs is the U_g coefficients, U_g is a sufficient cause of lethality.

Proposition 2 (network_forced is conservative). If $\text{verdict}(g) = \text{network_forced}$, then even an objective that demands none of the unproducible compounds cannot be met; the lesion is in shared precursor or energy supply. Proof: by (C1) applied to the rescued objective. Note that network_forced does not claim the gene is truly essential in vivo - only that no biomass relabelling rescues it.

Proposition 3 (minimality). U_g is minimal among sets defined by single-compound producibility: removing any i in U_g from the strip set leaves a compound with $\pi_i = 0$ in the objective, and a biomass reaction consuming an unproducible compound at positive coefficient forces $z^* = 0$. Hence every element of U_g is necessary for rescue.

Falsifiable prediction. If biomass_forced labels reflect model artifacts rather than biology, genes so labelled should be enriched for experimentally non-essential genes relative to network_forced genes. This is tested with the one-sided Fisher exact test (C18) in Section 4.4.

C.5 AUROC as a Mann-Whitney statistic

For scores s , positives P (essential, n_1) and negatives N (n_0):

$$AUROC = (1/(n_1 n_0)) \sum_{i \in P} \sum_{j \in N} [1(s_i > s_j) + \frac{1}{2} 1(s_i = s_j)] \quad (C6)$$

Derivation: the ROC curve is $(FPR(t), TPR(t))$ as t sweeps; integrating $TPR \, dFPR$ over the step function gives, for each negative j , the fraction of positives scoring above it, which summed over j and normalized is (C6). Ties receive half credit because a tied pair lies on a diagonal ROC segment. Consequence for FBA: binary FBA scores have massive ties, so AUROC collapses to $(TPR + TNR)/2$ at the single threshold, which explains why pure FBA AUROCs cluster near 0.63-0.67.

$$AUROC_{binary} = \frac{1}{2}(TPR + 1 - FPR) \quad (C7)$$

C.6 Average precision

$$AP = \sum_k (R_k - R_{k-1}) P_k \quad (C8)$$

with P_k, R_k precision and recall at the k -th threshold. The baseline for a random ranker equals the prevalence π (0.172 here), so AUPRC must always be read against π . The lift AP/π is reported alongside.

C.7 Matthews correlation coefficient

$$MCC = (TP \, TN - FP \, FN) / \sqrt{((TP+FP)(TP+FN)(TN+FP)(TN+FN))} \quad (C9)$$

MCC equals the Pearson correlation of the two binary indicator vectors (proof: expand the covariance of Bernoulli variables), so it is symmetric in classes and insensitive to prevalence in the sense that a constant predictor scores 0.

C.8 Percentile bootstrap confidence intervals

For statistic T and B resamples of genes with replacement, $T_b^* = T(X_b^*)$. The 95% interval is

$$CI_{95} = [Q_{0.025}(T^*), Q_{0.975}(T^*)] \quad (C10)$$

Validity rests on the bootstrap distribution of $T^* - T$ approximating that of $T - \theta$; for U-statistics such as AUROC this holds asymptotically (Hall, 1992). Resamples with a single class are discarded (the code reports `n_boot_used`). Genes are the resampling unit, so the interval reflects gene-set uncertainty, not experimental replicate noise.

C.9 Paired comparison of two classifiers: McNemar

On paired hard calls, let b = number of genes model A gets right and B wrong, c the reverse. Under H_0 (equal error rates) each discordant gene is a fair coin, $b \sim \text{Bin}(b+c, 1/2)$:

$$p = 2 \sum_{k=0}^{\min(b,c)} C(b+c, k) 2^{-(b+c)} \quad (C11)$$

$$\chi^2 = (|b - c| - 1)^2 / (b + c) \quad (C12)$$

(C12) is the continuity-corrected normal approximation, derived from the variance $(b+c)/4$ of $\text{Bin}(b+c, 1/2)$. With $b = 27$, $c = 50$ the exact test gives $p = 0.012$ against the ensemble at the 0.5 cut, which is why the paper reports the ranking gain and the hard-call loss together.

C.10 DeLong-free AUROC difference by paired bootstrap

$$\Delta = AUROC_A - AUROC_B, \quad CI \text{ from paired resamples (same gene indices for A and B)} \quad (C13)$$

Pairing preserves the correlation between models scored on identical genes and is strictly tighter than comparing two independent intervals.

C.11 Stacked generalization without leakage

Base scores are produced out-of-fold: for fold k with held-out set I_k , model $f^{(-k)}$ is trained on the complement and scores I_k . The meta-learner is

$$h(x) = \sigma(w_0 + \sum_m w_m f_m^{(-k(x))}(x)) \quad (C14)$$

Because each base prediction for gene x was made without seeing x, the meta-features are distributed as test-time features, and the meta-learner's own OOF evaluation is unbiased up to fold-to-fold variance (Wolpert, 1992). Non-learned FBA scores need no folding.

C.12 Logistic loss and its gradient

$$L(w) = -\sum_i [y_i \log p_i + (1-y_i) \log(1-p_i)] + \lambda ||w||^2/2, \quad \nabla L = X^T(p - y) + \lambda w \quad (C15)$$

Derivation: $d/dz \log \sigma(z) = 1 - \sigma(z)$. The Hessian $X^T \text{diag}(p(1-p))X + \lambda I$ is positive definite, so the meta-learner has a unique optimum.

C.13 Class-weighted loss for the CNN and GCN

$$L_w = -\sum_i [\omega_1 y_i \log p_i + \omega_0 (1-y_i) \log(1-p_i)], \quad \omega_c = n/(2 n_c) \quad (C16)$$

With $\omega_c = n/(2n_c)$ each class contributes equal total weight, which is the minimizer of expected balanced error for a calibrated model.

C.14 Receptive field of the sequence CNN

$$RF_L = 1 + \sum_{l=1}^L (k_l - 1) \prod_{j<l} s_j \quad (C17)$$

For kernel sizes k_l and strides s_j (pooling counted as stride). The receptive field bounds the longest motif the network can represent before global pooling; global max pooling then makes the gene score invariant to motif position.

C.15 Fisher exact test for provenance enrichment

Let the 2 x 2 table be [[a, b], [c, d]] with rows biomass_forced / network_forced and columns non-essential / essential in vivo. Conditioning on margins, a follows the hypergeometric law

$$P(A = a) = C(a+b, a) C(c+d, c) / C(n, a+c) \quad (C18)$$

$$P_{one-sided} = \sum_{k \geq a} P(A = k), \quad OR = ad / bc \quad (C19)$$

The one-sided alternative encodes the prediction of C.4: biomass_forced calls are more often wrong (non-essential in vivo) than network_forced calls.

C.16 Graph convolution as normalized propagation

$$H^{(l+1)} = \sigma(D^{-1/2} A D^{-1/2} H^{(l)} W^{(l)}), \quad A \blacksquare = A + I \quad (C20)$$

The symmetric normalization has spectrum in (-1, 1], so repeated propagation is contractive and does not explode; two layers aggregate the 2-hop gene neighbourhood via shared metabolites.

C.17 Dynamic FBA

$$dX/dt = \mu(t) X, \quad dS_k/dt = v_k(t) X, \quad \mu(t) = \max c^T v \text{ s.t. } Sv = 0, v_k \geq -V_{max,k} S_k / (K_{m,k} + S_k) \quad (C21)$$

$$X_{t+\Delta t} = X_t e^{\mu \Delta t}, \quad S_{t+\Delta t} = S_t + v (X_t / \mu) (e^{\mu \Delta t} - 1) \quad (C22)$$

(C22) is the exact solution of (C21) with fluxes frozen over the step (the static-optimization approach of Mahadevan et al., 2002); it avoids the negative-concentration artifacts of forward Euler. The Boolean catabolite-repression overlay sets $V_{max,lac} = 0$ while $S_{glc} > \theta$, producing diauxie.

C.18 Threshold sensitivity and bimodality

$$E(\tau) = 1(r_g < \tau), \quad \partial|E(\tau)|/\partial\tau = \sum_g \delta(r_g - \tau) \quad (C23)$$

Because FBA growth fractions are almost all exactly 0 or 1, the density of r_g in (0.001, 0.5) is nearly empty and $|E(\tau)|$ is flat across two orders of magnitude of τ - the robustness result of Section 3.

C.19 Lift and expected calibration

$$lift = AP / \pi, \quad ECE = \sum_b (|B_b|/n) |acc(B_b) - conf(B_b)| \quad (C24)$$

Lift normalizes AUPRC for prevalence, which differs between Gerdes (17.2%) and the PEC subset, making cross-yardstick comparisons meaningful.

Appendix D. Extended evaluation: distributions, calibration and early precision

AUROC summarizes the whole ranking, but a biologist using a virtual cell reads it from the top: which genes should be tested first? This section reports the views of the same out-of-fold scores that answer that question. Every panel is generated by `scripts/make_figures_ext.py` directly from `results/aligned_predictions_all_models.csv` and `results/ensemble_oof.csv`; no number is transcribed by hand.



Figure 6. Out-of-fold score distributions by Gerdes class for the GCN, sequence CNN, k-mer baseline and stacked ensemble. Only the ensemble develops a distinct high-score essential mode (scores > 0.8).

Figure 6 shows why the ensemble beats every member on ranking. Each learned base model shifts the essential class to the right but leaves the two classes overlapping across the full score range. The ensemble, which also receives the binary FBA calls, splits the essential genes into two groups: a sharp mode near 0.85-0.95 made of genes that FBA calls essential AND the learned models score high, and a broad low mode of essential genes that FBA misses. The second group is the ceiling of this architecture - no amount of stacking recovers genes that none of the inputs can see.




Figure 10. Precision among the top-k ranked genes. Dotted line: prevalence (0.172).

model	P@25	P@50	P@100	P@200
fba_min	0.32	0.44	0.46	0.42
gnn	0.52	0.44	0.44	0.34
cnn	0.12	0.32	0.34	0.28
kmer	0.44	0.34	0.38	0.33
ensemble	0.60	0.72	0.61	0.46

Table 8. Precision at k (fraction of the top- k genes that are essential in Gerdes 2003).

At $k = 50$ the ensemble reaches precision 0.72, 4.2x the prevalence and well above the best single model (FBA minimal 0.44, GCN 0.44). The advantage shrinks with k : by $k = 200$ the ensemble (0.46) is only marginally ahead of FBA (0.42). The practical reading is that the ensemble is a good triage tool for the first 50-100 experiments and not a replacement for FBA deeper in the list. The early dip in the ensemble curve at $k = 10$ is real and explained in Section 4: the very top of the ranking is occupied by molybdopterin genes (*moaD*, *moaE*, *moaA*, *moaC*) that every component scores as essential and that Gerdes finds dispensable - the biomass-objective artifact analysed in Section 4.3.



Figure 7. Reliability diagram (quantile bins; scores min-max scaled per model).

Figure 7 shows that none of the scores is calibrated as a probability: all curves lie below the diagonal, meaning that scaled scores overstate essentiality. This is expected for class-weighted training (Appendix C, eq. C16), which trades calibration for balanced error. The ensemble curve is the closest to monotone and reaches an observed essential fraction of about 0.55 in its top decile. For downstream use as probabilities, a Platt or isotonic recalibration on a held-out split is required; we do not claim calibrated probabilities anywhere in this paper.



Figure 8. Per-fold AUROC for the GCN and the sequence CNN (3-fold stratified CV).

Fold-to-fold spread (Figure 8) is about 0.05 AUROC for the GCN and 0.02 for the CNN. The GCN spread is comparable to the gap between the GCN and FBA, so the ordering GCN > FBA-rich in Table 1 should be read as a tie within fold noise. The ensemble gain (+0.057 over FBA minimal) is larger than both spreads and survives the paired bootstrap of Appendix C.10.



Figure 9. F1 and number of predicted essential genes versus the growth-fraction threshold.

Figure 9 is the graphical form of the bimodality result (eq. C23): both curves are flat, since exactly 183 genes are called essential at 0.5%, 1% and 5% of wild-type growth.

Table 9 lists the 40 highest-ranked genes so that any reader can check the ranking against EcoCyc or their own data. False positives cluster in cofactor biosynthesis (molybdopterin, folate), which motivated the provenance audit of Section 4.4.

rank	b-number	gene	ensemble	FBA min	GCN	Gerdes
1	b0784	moaD	0.963	E	0.81	non-ess.
2	b0785	moaE	0.958	E	0.76	non-ess.

rank	b-number	gene	ensemble	FBA min	GCN	Gerdess
3	b2746	ispF	0.952	E	0.72	essential
4	b0142	folK	0.948	E	0.81	essential
5	b0009	mog	0.944	E	0.71	essential
6	b3857	mobA	0.944	E	0.82	non-ess.
7	b3018	plsC	0.943	E	0.80	essential
8	b0783	moaC	0.939	E	0.75	non-ess.
9	b3058	folB	0.938	E	0.72	non-ess.
10	b3177	folP	0.934	E	0.79	non-ess.
11	b3634	coaD	0.930	E	0.67	essential
12	b2747	ispD	0.928	E	0.76	essential
13	b3807	cyaY	0.927	E	0.71	non-ess.
14	b0415	ribH	0.925	E	0.71	essential
15	b0029	ispH*	0.925	E	0.69	essential
16	b0159	mtn	0.922	E	0.48	essential
17	b0781	moaA	0.921	E	0.70	non-ess.
18	b3041	ribB	0.920	E	0.49	essential
19	b0369	hemB	0.919	E	0.58	essential
20	b1096	pabC	0.919	E	0.59	non-ess.
21	b1208	ispE	0.918	E	0.71	essential
22	b0826	moeB	0.917	E	0.67	non-ess.
23	b1277	ribA	0.917	E	0.44	essential
24	b0421	ispA	0.915	E	0.69	essential
25	b0174	uppS	0.913	E	0.54	essential
26	b3648	gmk	0.909	E	0.61	essential
27	b3639	dfp	0.908	E	0.60	essential
28	b3040	zupT	0.907	E	0.43	non-ess.
29	b1098	tmk	0.905	E	0.65	essential
30	b0103	coaE	0.904	E	0.70	essential
31	b1662	ribE	0.903	E	0.66	essential
32	b0154	hemL	0.902	E	0.54	essential
33	b2515	ispG*	0.901	E	0.73	essential
34	b1415	aldA	0.894	E	0.48	non-ess.
35	b0417	thiL	0.894	E	0.62	essential
36	b4039	ubiC	0.893	E	0.48	non-ess.
37	b0475	hemH	0.893	E	0.83	essential
38	b0087	mraY	0.890	E	0.56	non-ess.
39	b2530	iscS	0.890	E	0.53	essential
40	b2585	pssA	0.889	E	0.75	essential

Table 9. Top 40 genes by out-of-fold ensemble score with the Gerdess 2003 call.

Appendix E. Essentiality-provenance catalog for iJO1366

E.1 Question and design

Section 4.3 showed one case - the molybdopter cluster - where an FBA essentiality call is produced by the biomass objective rather than by the network. This appendix asks the general question: for every one of the 208 genes that iJO1366 calls essential on glucose minimal medium, is the call `network_forced` or `biomass_forced`? The rescue test (Appendix C.4, eqs. C3-C5) answers it per gene with one re-runnable LP experiment. The tool is shipped as `python -m vcell audit-biomass MODEL` and runs on any COBRA model file.

Prior work. The phenomenon of cofactor-driven false essentiality is known. Bernstein et al. (2023, Mol Syst Biol) found 21 vitamin/cofactor biosynthesis genes in iML1515 (biotin, R-pantothenate, thiamin, tetrahydrofolate, NAD+) that the model calls essential and RB-TnSeq finds fit, and fixed them by manually supplying those cofactors. MEMOTE tests whether biomass precursors are producible in the wild-type model but does not attribute any knockout's lethality to specific compounds. Our contribution is narrower and procedural: an automated, per-gene, falsifiable attribution. We found no existing tool that does this; that is a statement about our search, not a proof of absence.

E.2 Verdict census

verdict	genes
biomass_forced	208

Table 10. Rescue-test verdicts.

Negative result 1: the binary verdict does not discriminate. All 208 iJO1366 essential genes are `biomass_forced`. Every lethal knockout on glucose minimal medium makes at least one biomass constituent unproducible, and removing those constituents always restores some growth. Proposition 1 is true but uninformative here: the set U_g ranges from 1 to 50 compounds (median 2), and when U_g is large the rescued objective is a different, much easier objective. The rescue test therefore needs a second axis before it says anything about artifacts.

E.3 A methods catch: raw `biomass_forced` over-calls artifacts

A raw `biomass_forced` label means only that stripping the unproducible compounds rescues growth. For a gene such as `fabD` that stops lipid synthesis this is trivially true and says nothing about artifacts: lipids are really required. The artifact hypothesis only makes sense for compounds a mutant might plausibly obtain without making them (trace cofactors and vitamins by carry-over or cross-feeding, the class named by Bernstein et al.). We therefore split biomass constituents into cofactor/vitamin, inorganic and bulk classes (`vcell/compound_classes.py`) and label a gene `cofactor_only` when every unproducible compound is a cofactor or vitamin.

stratum	genes	Gerdes essential	fraction
biomass_forced:bulk_involved	104	39	0.38
biomass_forced:cofactor_only	79	42	0.53

Table 11. Gerdes 2003 essentiality by provenance stratum.

E.4 The falsifiable test

Prediction (Appendix C.4): cofactor_only biomass_forced genes are enriched for experimentally non-essential genes relative to all other FBA-essential genes. One-sided Fisher exact test (eqs. C18-C19) on $[[\text{non-ess}, \text{ess}], [\text{non-ess}, \text{ess}]] = [[37, 42], [65, 39]]$: odds ratio 0.53, $p = 0.988$. For the raw (unstratified) biomass_forced label the table is $[[102, 81], [0, 0]]$ (degenerate: the network_forced row is empty, so no test is possible): odds ratio nan, $p = 1$.

Negative result 2: the cofactor-artifact hypothesis fails on this yardstick. Cofactor-only genes are more often essential in Gerdes 2003 than bulk-precursor genes, the opposite of the prediction, and the one-sided test is nowhere near significant. The likely reason is a condition mismatch. Gerdes screened on rich LB medium, where amino acids and nucleotide precursors are supplied, so bulk-precursor biosynthesis genes are dispensable, while most cofactors are not supplied at useful levels, so cofactor genes stay essential. The model was simulated on glucose minimal medium. A condition-matched test (RB-TnSeq on defined carbon sources, Appendix H) is the right place to test the hypothesis. We keep the failed test in the paper because it shows that the yardstick has to match the condition before essentiality errors can be attributed.

E.5 Corrected benchmark

metric	FBA iJO1366 (original)	after cofactor_only correction
auroc	0.666	0.584
auprc	0.316	0.247
precision	0.437	0.369
recall	0.386	0.191
f1	0.410	0.252
mcc	0.297	0.163

Table 12. iJO1366 minimal-medium FBA vs Gerdes 2003 before and after flipping cofactor_only biomass_forced genes to non-essential.

Corrected AUROC 95% CI: [0.544, 0.626]. The correction uses only the audit, never the Gerdes labels, and it makes the model worse on this yardstick, consistent with negative result 2. It is not adopted.

biomass compound	genes whose knockout makes it unproducible
btn_c	37
bmocogdp_c	28
thmpp_c	28
coa_c	27
10fthf_c	26
mlthf_c	26
thf_c	26
amet_c	23

biomass compound	genes whose knockout makes it unproducible
murein5px4p_p	22
fad_c	19
nadp_c	17
kdo2lipid4_e	17
ribflv_c	17
2ohph_c	16
nad_c	16
dctp_c	12
gtp_c	12
met__L_c	12
atp_c	11
datp_c	11
his__L_c	11
pheme_c	10
sheme_c	10
udcpdp_c	9
trp__L_c	9
lys__L_c	8
arg__L_c	8
phe__L_c	8
2fe2s_c	8
4fe4s_c	8
leu__L_c	7
tyr__L_c	7
dttp_c	7
cys__L_c	7
ile__L_c	6
ctp_c	6
dctp_c	6
pydx5p_c	5
pe160_p	5
pe161_p	5
utp_c	5
pe160_c	4
pe161_c	4
thr__L_c	3
pro__L_c	3
gln__L_c	2
glu__L_c	2

biomass compound	genes whose knockout makes it unproducible
val__L_c	2
asn__L_c	1
asp__L_c	1
cobalt2_c	1
cu2_c	1
ca2_c	1

Table 13. Compound census over biomass_forced genes.

E.6 Full catalog

Every audited gene with its verdict, the unproducible biomass compounds (truncated to 60 characters) and its Gerdes 2003 call (E essential, N non-essential, n/a not in the aligned set). * marks network_forced_no_biomass_link: no individual biomass constituent becomes unproducible, yet growth stops, which points to a coupling or energy constraint.

b-number	gene	verdict	unproducible compounds	Gerdes
b0003		biomass_forced	ile__L_c, thr__L_c	n/a
b0004	thrC	biomass_forced	ile__L_c, thr__L_c	N
b0009	mog	biomass_forced	bmocogdp_c	E
b0025	ribF	biomass_forced	fad_c	E
b0029	ispH*	biomass_forced	2ohph_c, ucdpdp_c	E
b0031	dapB	biomass_forced	lys__L_c, murein5px4p_p	E
b0052	pdxA	biomass_forced	pydx5p_c	N
b0071	leuD	biomass_forced	leu__L_c	N
b0072	leuC	biomass_forced	leu__L_c	N
b0073	leuB	biomass_forced	leu__L_c	N
b0074	leuA	biomass_forced	leu__L_c	N
b0085	murE	biomass_forced	murein5px4p_p	E
b0086	murF	biomass_forced	murein5px4p_p	N
b0087	mraY	biomass_forced	murein5px4p_p	N
b0088	murD	biomass_forced	murein5px4p_p	E
b0090	murG	biomass_forced	murein5px4p_p	E
b0091	murC	biomass_forced	murein5px4p_p	E
b0096		biomass_forced	kdo2lipid4_e	n/a
b0103	coaE	biomass_forced	coa_c	E
b0109	nadC	biomass_forced	nad_c, nadp_c	N
b0131	panD	biomass_forced	coa_c	N
b0133	panC	biomass_forced	coa_c	N
b0134	panB	biomass_forced	coa_c	N
b0142	folK	biomass_forced	10fthf_c, mlthf_c, thf_c	E
b0154	hemL	biomass_forced	pheme_c, sheme_c	E

b-number	gene	verdict	unproducible compounds	Gerdes
b0159	mtn	biomass_forced	btn_c, sheme_c, thmpp_c	E
b0166	dapD	biomass_forced	lys__L_c, murein5px4p_p	E
b0173	dxr	biomass_forced	2ohph_c, ucdpdp_c	E
b0174	uppS	biomass_forced	ucdpdp_c	E
b0175	cdsA	biomass_forced	pe160_c, pe160_p, pe161_c, pe161_p	E
b0179		biomass_forced	kdo2lipid4_e	n/a
b0180	fabZ	biomass_forced	btn_c	E
b0181	lpxA	biomass_forced	kdo2lipid4_e	E
b0182	lpxB	biomass_forced	kdo2lipid4_e	E
b0185	accA	biomass_forced	btn_c, kdo2lipid4_e	E
b0369	hemB	biomass_forced	pheme_c, sheme_c	E
b0386	proC	biomass_forced	pro__L_c	E
b0414	ribD	biomass_forced	fad_c, ribflv_c	E
b0415	ribH	biomass_forced	fad_c, ribflv_c	E
b0417	thiL	biomass_forced	thmpp_c	E
b0420	dxs	biomass_forced	2ohph_c, pydx5p_c, thmpp_c, ucdpdp_c	E
b0421	ispA	biomass_forced	2ohph_c, ucdpdp_c	E
b0423	thiI	biomass_forced	thmpp_c	N
b0475	hemH	biomass_forced	pheme_c	E
b0522	purK	biomass_forced	10fthf_c, amet_c, atp_c, bmocogdp_c, btn_c, coa_c, datp_c, d	N
b0523	purE	biomass_forced	10fthf_c, amet_c, atp_c, bmocogdp_c, btn_c, coa_c, datp_c, d	N
b0524	lpxH	biomass_forced	kdo2lipid4_e	E
b0639	nadD	biomass_forced	nad_c, nadp_c	N
b0720	gltA	biomass_forced	10fthf_c, arg__L_c, gln__L_c, glu__L_c, mlthf_c, murein5px4p	N
b0750	nadA	biomass_forced	nad_c, nadp_c	N
b0774	bioA	biomass_forced	btn_c	N
b0775	bioB	biomass_forced	btn_c	N
b0776	bioF	biomass_forced	btn_c	N
b0777	bioC	biomass_forced	btn_c	N
b0778	bioD	biomass_forced	btn_c	N
b0781	moaA	biomass_forced	bmocogdp_c	N
b0783	moaC	biomass_forced	bmocogdp_c	N
b0784	moaD	biomass_forced	bmocogdp_c	N
b0785	moaE	biomass_forced	bmocogdp_c	N
b0826	moeB	biomass_forced	bmocogdp_c	N
b0827	moeA	biomass_forced	bmocogdp_c	N
b0907	serC	biomass_forced	pydx5p_c	N
b0908	aroA	biomass_forced	10fthf_c, 2ohph_c, mlthf_c, phe__L_c, thf_c, thmpp_c, trp__L	N
b0914	msbA	biomass_forced	kdo2lipid4_e, pe160_p, pe161_p	E

b-number	gene	verdict	unproducible compounds	Gerdas
b0915	lpxK	biomass_forced	kdo2lipid4_e	E
b0918	kdsB	biomass_forced	kdo2lipid4_e	E
b0928	aspC	biomass_forced	10fthf_c, amet_c, arg__L_c, asn__L_c, asp__L_c, atp_c, bmoco	N
b1062	pyrC	biomass_forced	ctp_c, dctp_c, dttp_c, utp_c	N
b1069	mviN	biomass_forced	murein5px4p_p	E
b1091	fabH	biomass_forced	btn_c	E
b1092	fabD	biomass_forced	btn_c, kdo2lipid4_e	E
b1093	fabG	biomass_forced	btn_c, kdo2lipid4_e	E
b1094		biomass_forced	btn_c, kdo2lipid4_e, pe160_c, pe160_p, pe161_c, pe161_p	n/a
b1096	pabC	biomass_forced	10fthf_c, mlthf_c, thf_c	N
b1098	tmk	biomass_forced	dttp_c	E
b1131	purB	biomass_forced	10fthf_c, amet_c, atp_c, bmocogdp_c, btn_c, coa_c, datp_c, d	N
b1136	icd	biomass_forced	10fthf_c, arg__L_c, gln__L_c, glu__L_c, mlthf_c, murein5px4p	N
b1208	ispE	biomass_forced	2ohph_c, ucdpdp_c	E
b1210		biomass_forced	pheme_c, sheme_c	n/a
b1215		biomass_forced	kdo2lipid4_e	n/a
b1260	trpA	biomass_forced	trp__L_c	N
b1261	trpB	biomass_forced	trp__L_c	N
b1262	trpC	biomass_forced	trp__L_c	N
b1263	trpD	biomass_forced	trp__L_c	N
b1264	trpE	biomass_forced	trp__L_c	N
b1277	ribA	biomass_forced	fad_c, ribflv_c	E
b1281	pyrF	biomass_forced	ctp_c, dctp_c, dttp_c, utp_c	E
b1288	fabI	biomass_forced	btn_c	E
b1415	aldA	biomass_forced	10fthf_c, mlthf_c, thf_c	N
b1662	ribE	biomass_forced	fad_c, ribflv_c	E
b1693	aroD	biomass_forced	10fthf_c, 2ohph_c, mlthf_c, phe__L_c, thf_c, thmpp_c, trp__L	N
b1740	nadE	biomass_forced	nad_c, nadp_c	E
b1812	pabB	biomass_forced	10fthf_c, mlthf_c, thf_c	N
b2019		biomass_forced	his__L_c	n/a
b2020	hisD	biomass_forced	his__L_c	N
b2021	hisC	biomass_forced	his__L_c	N
b2022	hisB	biomass_forced	his__L_c	N
b2023	hisH	biomass_forced	his__L_c	N
b2024	hisA	biomass_forced	his__L_c	N
b2025	hisF	biomass_forced	his__L_c	N
b2026	hisI	biomass_forced	his__L_c	N
b2103	thiD	biomass_forced	thmpp_c	N
b2153	folE	biomass_forced	10fthf_c, mlthf_c, thf_c	E

b-number	gene	verdict	unproducible compounds	Gerdes
b2311	ubiX	biomass_forced	2ohph_c	E
b2312	purF	biomass_forced	10fthf_c, amet_c, atp_c, bmocogdp_c, btn_c, coa_c, datp_c, d	N
b2315		biomass_forced	10fthf_c, mlthf_c, thf_c	n/a
b2316		biomass_forced	btn_c, kdo2lipid4_e	n/a
b2320	pdxB	biomass_forced	pydx5p_c	N
b2323	fabB	biomass_forced	btn_c, kdo2lipid4_e	E
b2329	aroC	biomass_forced	10fthf_c, 2ohph_c, mlthf_c, phe__L_c, thf_c, thmpp_c, trp__L	N
b2400	gltX	biomass_forced	pheme_c, sheme_c	E
b2472	dapE	biomass_forced	lys__L_c, murein5px4p_p	E
b2476	purC	biomass_forced	10fthf_c, amet_c, atp_c, bmocogdp_c, btn_c, coa_c, datp_c, d	N
b2478	dapA	biomass_forced	lys__L_c, murein5px4p_p	E
b2499	purM	biomass_forced	10fthf_c, amet_c, atp_c, bmocogdp_c, btn_c, coa_c, datp_c, d	N
b2507	guaA	biomass_forced	10fthf_c, bmocogdp_c, dgtp_c, fad_c, gtp_c, mlthf_c, ribflv__	N
b2515	ispG*	biomass_forced	2ohph_c, ucdpdp_c	E
b2530	iscS	biomass_forced	bmocogdp_c, thmpp_c	E
b2557	purL	biomass_forced	10fthf_c, amet_c, atp_c, bmocogdp_c, btn_c, coa_c, datp_c, d	N
b2564	pdxJ	biomass_forced	pydx5p_c	E
b2574	nadB	biomass_forced	nad_c, nadp_c	N
b2585	pssA	biomass_forced	pe160_c, pe160_p, pe161_c, pe161_p	E
b2599	pheA	biomass_forced	phe__L_c, thmpp_c, tyr__L_c	N
b2600	tyrA	biomass_forced	phe__L_c, thmpp_c, tyr__L_c	N
b2615	ppnK	biomass_forced	nadp_c	E
b2687		biomass_forced	btn_c, sheme_c	n/a
b2746	ispF	biomass_forced	2ohph_c, ucdpdp_c	E
b2747	ispD	biomass_forced	2ohph_c, ucdpdp_c	E
b2750		biomass_forced	2fe2s_c, 4fe4s_c, amet_c, bmocogdp_c, btn_c, coa_c, cys__L_c	n/a
b2751	cysN	biomass_forced	2fe2s_c, 4fe4s_c, amet_c, bmocogdp_c, btn_c, coa_c, cys__L_c	N
b2752	cysD	biomass_forced	2fe2s_c, 4fe4s_c, amet_c, bmocogdp_c, btn_c, coa_c, cys__L_c	N
b2762	cysH	biomass_forced	2fe2s_c, 4fe4s_c, amet_c, bmocogdp_c, btn_c, coa_c, cys__L_c	N
b2763	cysI	biomass_forced	2fe2s_c, 4fe4s_c, amet_c, bmocogdp_c, btn_c, coa_c, cys__L_c	N
b2764	cysJ	biomass_forced	2fe2s_c, 4fe4s_c, amet_c, bmocogdp_c, btn_c, coa_c, cys__L_c	N
b2780	pyrG	biomass_forced	ctp_c, dctp_c	E
b2818	argA	biomass_forced	arg__L_c	N
b2827	thyA	biomass_forced	dttp_c	E
b2838	lysA	biomass_forced	lys__L_c	N
b2942	metK	biomass_forced	amet_c, btn_c, sheme_c, thmpp_c	E
b3018	plsC	biomass_forced	pe160_c, pe160_p, pe161_c, pe161_p	E
b3040	zupT	biomass_forced	cobalt2_c, cu2_c	N
b3041	ribB	biomass_forced	fad_c, ribflv_c	E

b-number	gene	verdict	unproducible compounds	Gerdas
b3058	folB	biomass_forced	10fthf_c, mlthf_c, thf_c	N
b3172	argG	biomass_forced	arg__L_c	N
b3176	mrsA	biomass_forced	kdo2lipid4_e, murein5px4p_p	N
b3177	folP	biomass_forced	10fthf_c, mlthf_c, thf_c	N
b3187	ispB	biomass_forced	2ohph_c	E
b3189	murA	biomass_forced	murein5px4p_p	N
b3196	yrbG	biomass_forced	ca2_c	N
b3198		biomass_forced	kdo2lipid4_e	n/a
b3199		biomass_forced	kdo2lipid4_e	n/a
b3200		biomass_forced	kdo2lipid4_e	n/a
b3201	yhbG	biomass_forced	kdo2lipid4_e	E
b3255		biomass_forced	btn_c, kdo2lipid4_e	n/a
b3256		biomass_forced	btn_c, kdo2lipid4_e	n/a
b3359		biomass_forced	arg__L_c, lys__L_c, murein5px4p_p	n/a
b3360	pabA	biomass_forced	10fthf_c, mlthf_c, thf_c	N
b3368		biomass_forced	sheme_c	n/a
b3389	aroB	biomass_forced	10fthf_c, 2ohph_c, mlthf_c, phe__L_c, thf_c, thmpp_c, trp__L	N
b3412	bioH	biomass_forced	btn_c	N
b3433	asd	biomass_forced	amet_c, btn_c, ile__L_c, lys__L_c, met__L_c, murein5px4p_p,	E
b3607	cysE	biomass_forced	2fe2s_c, 4fe4s_c, amet_c, bmocogdp_c, btn_c, coa_c, cys__L_c	E
b3608	gpsA	biomass_forced	pe160_c, pe160_p, pe161_c, pe161_p	N
b3633	kdtA	biomass_forced	kdo2lipid4_e	E
b3634	coaD	biomass_forced	coa_c	E
b3639	dfp	biomass_forced	coa_c	E
b3642	pyrE	biomass_forced	ctp_c, dctp_c, dttp_c, utp_c	N
b3648	gmK	biomass_forced	10fthf_c, bmocogdp_c, dgtp_c, fad_c, gtp_c, mlthf_c, ribflv__	E
b3729	glmS	biomass_forced	kdo2lipid4_e, murein5px4p_p	E
b3730	glmU	biomass_forced	kdo2lipid4_e, murein5px4p_p	E
b3770	ilvE	biomass_forced	ile__L_c, leu__L_c, phe__L_c	N
b3771	ilvD	biomass_forced	coa_c, ile__L_c, leu__L_c, val__L_c	N
b3774	ilvC	biomass_forced	coa_c, ile__L_c, leu__L_c, val__L_c	N
b3804	hemD	biomass_forced	pheme_c, sheme_c	E
b3805	hemC	biomass_forced	pheme_c, sheme_c	E
b3807	cyaY	biomass_forced	2fe2s_c, 4fe4s_c	N
b3809	dapF	biomass_forced	lys__L_c, murein5px4p_p	E
b3843	ubiD	biomass_forced	2ohph_c	N
b3850	hemG	biomass_forced	pheme_c	E
b3857	mobA	biomass_forced	bmocogdp_c	N
b3939	metB	biomass_forced	amet_c, btn_c, met__L_c	N

b-number	gene	verdict	unproducible compounds	Gerdas
b3941	metF	biomass_forced	amet_c, btn_c, met__L_c, sheme_c	N
b3957		biomass_forced	arg__L_c	n/a
b3958	argC	biomass_forced	arg__L_c	E
b3959	argB	biomass_forced	arg__L_c	N
b3960	argH	biomass_forced	arg__L_c	N
b3967	murI	biomass_forced	murein5px4p_p	E
b3972	murB	biomass_forced	murein5px4p_p	N
b3974	coaA	biomass_forced	coa_c	E
b3990		biomass_forced	thmpp_c	n/a
b3991		biomass_forced	thmpp_c	n/a
b3992	thiF	biomass_forced	thmpp_c	N
b3993	thiE	biomass_forced	thmpp_c	N
b3994	thiC	biomass_forced	thmpp_c	N
b3997	hemE	biomass_forced	pheme_c	E
b4005	purD	biomass_forced	10fthf_c, amet_c, atp_c, bmocogdp_c, btn_c, coa_c, datp_c, d	N
b4006	purH	biomass_forced	10fthf_c, amet_c, atp_c, bmocogdp_c, btn_c, coa_c, datp_c, d	N
b4013	metA	biomass_forced	amet_c, btn_c, met__L_c	N
b4039	ubiC	biomass_forced	2ohph_c	N
b4040		biomass_forced	2ohph_c	n/a
b4160		biomass_forced	pe160_c, pe160_p, pe161_c, pe161_p	n/a
b4177	purA	biomass_forced	amet_c, atp_c, btn_c, coa_c, datp_c, fad_c, his__L_c, nad_c,	N
b4214	cysQ	biomass_forced	2fe2s_c, 4fe4s_c, amet_c, bmocogdp_c, btn_c, coa_c, cys__L_c	N
b4245	pyrB	biomass_forced	ctp_c, dctp_c, dttp_c, utp_c	N
b4261	yjgP	biomass_forced	kdo2lipid4_e	E
b4262	yjgQ	biomass_forced	kdo2lipid4_e	E
b4407		biomass_forced	thmpp_c	n/a
s0001		biomass_forced	10fthf_c, ala__L_c, amet_c, arg__L_c, asn__L_c, asp__L_c, at	n/a

Table 14. Essentiality-provenance catalog.

Appendix F. External evidence and ensemble v2

F.1 Motivation

Figure 6 showed that the v1 ensemble cannot recover essential genes that none of its inputs sees. The obvious remedy is new, independent evidence about each gene that uses no essentiality labels: its position in the protein-association network, its physical properties and its pathway context. We added three external resources, each snapshotted under `data/external/` and processed by `vcell/external_features.py`: STRING v12 (organism 511145), the UniProt reference proteome UP000000625, and KEGG REST gene-pathway links.

F.2 Features

From STRING we compute, per gene, total degree and weighted degree over all links, and on the high-confidence subgraph (score ≥ 700) degree, k-core number, local clustering coefficient and PageRank (networkx). PageRank is the stationary distribution of

$$PR(i) = (1 - d)/N + d \sum_{j \in N(i)} PR(j) / \deg(j), \quad d = 0.85 \text{ (F1)}$$

Because STRING's experimental, database and text-mining channels may encode prior knowledge about essentiality (well-studied essential genes accumulate literature links), we also rebuild the network from the genomic-context and coexpression channels only, combining channel scores s_c as

$$s = 1 - \prod_{c \in \{\text{neighborhood}, \text{fusion}, \text{cooccurrence}, \text{coexpression}\}} (1 - s_c / 1000) \text{ (F2)}$$

which is STRING's noisy-OR combination without its prior correction. Features from this network carry the prefix `strctx`. From UniProt we take log protein length, transmembrane segment count, presence of a cofactor annotation, number of EC numbers, annotation score and cytoplasm/membrane localization. From KEGG we take the number of pathways and membership in a fixed list of 24 core pathways (central carbon, nucleotide, amino-acid, cofactor, cell-wall, ribosome, replication, RNA polymerase, aminoacyl-tRNA, protein export).

F.3 Univariate associations

Each feature was tested against the Gerdes labels with a two-sided Mann-Whitney U test and the p-values corrected with Benjamini-Hochberg (statsmodels):

$$q_{(k)} = \min_{j \geq k} \min(1, m p_{(j)} / j) \text{ (F3)}$$

feature	AUROC	median ess.	median non-ess.	p	q (BH)
strctx_wdeg_all	0.687	57.3	38.8	5e-18	6.2e-17
strctx_deg_all	0.687	279	188	5.1e-18	6.2e-17
string_deg_all	0.676	334	243	3.6e-16	2.9e-15
up_cytoplasm	0.624	0	0	5.2e-15	3.1e-14
kegg_core	0.632	1	0	4.7e-14	2.2e-13
pax_log_ppm	0.639	2.06	1.44	1.5e-10	5.4e-10
string_wdeg_all	0.639	102	71	1.6e-10	5.4e-10

feature	AUROC	median ess.	median non-ess.	p	q (BH)
up_len	0.383	5.73	5.9	6.7e-08	2e-07
up_membrane	0.432	0	0	0.00012	0.00031
up_tm	0.437	0	0	0.00013	0.00031
string_core	0.581	11	10	0.00018	0.00038
kegg_n_path	0.578	3	2	0.00028	0.00055
string_deg_hi	0.572	19	16	0.00081	0.0015
up_n_ec	0.556	1	1	0.0024	0.0042
gc3	0.444	0.569	0.582	0.0094	0.015
strctx_deg_hi	0.552	6	5	0.015	0.023
strctx_clust	0.451	0.333	0.429	0.022	0.031
up_annot	0.540	5	5	0.031	0.042
cai	0.543	0.61	0.606	0.045	0.057
string_pagerank	0.531	2.66	2.6	0.15	0.18
strctx_core	0.530	4	4	0.17	0.2
string_clust	0.473	0.405	0.421	0.21	0.23
strctx_pagerank	0.522	2.82	2.83	0.31	0.32
up_cofactor	0.490	0	0	0.59	0.59

Table 15. Univariate association of external features with Gerdes 2003 essentiality.

Finding. The single strongest external predictor is the degree in the leakage-controlled genomic-context/coexpression network (strctx_wdeg_all, AUROC 0.687, $q = 6e-17$). It is stronger than the degree in STRING's full combined network, which includes literature and database evidence. The signal that essential genes are network hubs therefore does not come from literature bias. Essential proteins are also shorter than non-essential ones (AUROC 0.383 for length, i.e. 0.617 for shortness), and less often membrane proteins.

F.4 Ensemble v2 and an honest decomposition of the gain

model	OOF AUROC	OOF AUPRC
v1 as published (unscaled stack)	0.722	0.414
v1 re-fit with feature scaling	0.764	0.450
external features only (LR)	0.750	0.424
v2 LR (pre-specified primary)	0.797	0.498
v2 gradient boosting	0.777	0.532
v2 rank-average LR+GBM	0.800	0.543
v2 leakage-controlled LR	0.788	0.485
v2 leakage-controlled GBM	0.786	0.546
v2 leakage-controlled rank-average	0.803	0.545

Table 16. Out-of-fold performance vs Gerdes 2003 (1,249 genes, 3-fold stratified CV, seed 7).

The improvement over the published v1 number has two sources, and we report them separately. First, a defect in v1: the k-mer scores span only 0.46-0.54, so an unscaled logistic stack gives them almost

no weight. Standardizing features before stacking lifts v1 from 0.722 to 0.764 with no new information. Second, the external evidence: v2 LR improves on the re-fit v1 by +0.033 AUROC (paired bootstrap 95% CI [0.009, 0.058], one-sided $p = 0.002$). With the leakage-controlled feature set the gain is +0.024 [0.001, 0.047]. That is smaller but still above zero. The primary model was fixed as LR before seeing results. The rank-average models score slightly higher but were chosen afterwards and are reported as secondary. Averaged over five CV seeds, v2 LR scores 0.794.

A further feature block (v3) was tested and did not help: PaxDb integrated protein abundance and the codon adaptation index (CAI, Sharp and Li 1987; $w_c = f_c / \max_{c' \in \text{syn}(c)} f_{c'}$, CAI = $\exp(\text{mean log } w)$ against the ribosomal-protein reference set) with GC3. Abundance alone is associated with essentiality (Table 15), but adding it changes v2 LR by -0.002 AUROC [-0.009, 0.004]. Its signal is already captured by network degree and localization. v2 stays the primary model.

What this is not: it is not a state-of-the-art claim. Published machine-learning essentiality predictors for *E. coli* use different gene sets, labels and splits, so we have no like-for-like number to beat. The claim is only the internal, paired comparison above.

Appendix G. External tools and dataset manifest

This appendix lists every external tool (35) and every accession-level dataset (23) actually used in this work. A condition matrix from one study counts as one dataset (the Keio_ML9 RB-TnSeq matrix has 92 conditions but is listed once). Nothing is listed that was only considered.

#	tool	kind	used for
1	COBRApy 0.32	package	FBA, gene-deletion screens, dFBA, provenance audit
2	optlang 1.9	package	LP modelling layer under COBRApy
3	GLPK (swiglpk 5.0)	solver	linear-programming solver for every FBA call
4	python-libsbml 5.21	package	reading Bernstein SBML models (iML1515, all-corrections model)
5	PyTorch 2.x (CPU)	package	1D sequence CNN and pure-torch GCN
6	NumPy	package	array maths throughout
7	pandas	package	data alignment, joins, result tables
8	SciPy	package	Fisher exact, chi-square/McNemar, statistics
9	scikit-learn	package	logistic stacking, gradient boosting, CV, AUROC/AUPRC, calibration
10	networkx	package	STRING graph centralities (core number, clustering, PageRank)
11	statsmodels	package	multiple-testing correction (Benjamini-Hochberg)
12	Biopython	package	GenBank CDS extraction for U00096.3
13	matplotlib	package	all figures
14	Pillow	package	figure sizing and visual QA composites
15	ReportLab	package	Times-Roman PDF typesetting
16	pypdf	package	page-count verification
17	poppler pdffonts	CLI	embedded-font verification
18	pytest	package	hermetic test suite
19	git	CLI	version control, bundles
20	BiGG Models	database	e_coli_core and iJO1366 model downloads
21	NCBI Entrez E-utilities	web API	U00096.3 genome FASTA and GenBank
22	genome.wisc.edu Gerdes 2003 supplement	web resource	essentiality labels (Tables S1, S2)
23	EcoliWiki	web resource	essential-gene list cross-check
24	LBL Fitness Browser / RB-TnSeq downloads	web resource	Keio_ML9 fitness matrix and experiment metadata
25	Bernstein et al. 2023 pipeline (GitHub dbernste/E_coli_GEM_validation)	published tool	head-to-head benchmark: loading, matching, adjustments, simulation, PR-AUC metric
26	STRING v12	database	E. coli 511145 protein links (combined + per-channel)
27	UniProt REST	web API	proteome UP000000625 protein features
28	KEGG REST	web API	eco gene-pathway links and pathway names
29	PubMed Central	literature	novelty check, benchmark protocols
30	MEMOTE documentation	tool docs	novelty check of biomass tests
31	GitHub	platform	repository hosting, Bernstein code
32	Google Drive	platform	paper and results delivery

#	tool	kind	used for
33	PaxDb	database	integrated whole-organism protein abundance (v3 feature test)
34	ESM-2 t6-8M (fair-esm 2.0; Lin et al. 2023)	deep-learning model	protein language-model embeddings for all 4,300 CDS; v4 essentiality feature test (negative)
35	goatools 1.6 (GOEnrichmentStudy)	package	GO enrichment of ensemble hard errors (Fisher + BH)

Table 17. External tools.

#	dataset	type
1	BiGG e_coli_core	metabolic model
2	BiGG iJO1366	metabolic model
3	iML1515 (Bernstein repo copy of BiGG)	metabolic model
4	iJR904 (Bernstein repo)	metabolic model
5	iAF1260 (Bernstein repo)	metabolic model
6	Bernstein iML1515 all-corrections model	metabolic model
7	NCBI U00096.3 FASTA	genome
8	NCBI U00096.3 GenBank	genome annotation
9	Gerdes 2003 Table S1	essentiality screen
10	Gerdes 2003 Table S2 cross-study consensus	essentiality consensus
11	EcoliWiki essential genes	essentiality list
12	Price et al. 2018 Keio_ML9 fit_logratios_good (92 conditions)	RB-TnSeq fitness
13	Keio_ML9 expsUsed metadata	experiment metadata
14	Bernstein fit_organism_Keio + mapped media/carbon tables	curated fitness + media mapping
15	STRING v12 511145 protein.links	PPI network
16	STRING v12 511145 protein.links.detailed	PPI evidence channels
17	STRING v12 511145 protein.info	protein annotation
18	UniProt UP000000625	proteome
19	KEGG eco gene-pathway links	pathway map
20	KEGG eco pathway list	pathway names
21	PaxDb 511145 WHOLE_ORGANISM integrated	protein abundance
22	Gene Ontology go-basic.obo (current release)	ontology
23	GO Consortium EcoCyc GAF (taxon 83333)	GO annotations

Table 18. Dataset manifest.

Appendix H. Head-to-head on the Bernstein et al. (2023) RB-TnSeq benchmark

H.1 Why this benchmark

Appendix E showed that Gerdes 2003 (rich LB medium) cannot test whether cofactor requirements in the biomass objective create false essentiality calls, because the model and the experiment use different media. Bernstein et al. (2023) built the condition-matched benchmark for exactly this question. It compares FBA growth calls with RB-TnSeq fitness (Price et al. 2018, E. coli BW25113 Keio_ML9 library) for every model gene on 25 defined carbon sources. They report that hand-supplying five vitamins (biotin, R-pantothenate, thiamin, tetrahydrofolate, NAD+) fixes many false negatives in iML1515, and they release an all-corrections model. Their fully corrected analysis (all-corrections model plus the five vitamins) is reported at PR-AUC 0.843 (0.844 with transaldolase constrained); that is the published state of the art for this task. The all-corrections model file alone, without vitamins, scores lower in their pipeline.

H.2 Protocol

To rule out any protocol drift, we do not re-implement their benchmark. scripts/run_bernstein_benchmark.py executes the function cells of their released notebook (MIT licence, vendored under external/E_coli_GEM_validation) for model loading, gene and carbon-source matching, the BW25113 strain correction, removal of non-conditionally essential genes, carbon-source mapping, FBA simulation (carbon uptake -10, other media -1000 mmol/gDW/h, growth threshold 0.001) and the metric. We changed two things: their model_adjustments function reads a notebook-global variable, which we bind explicitly; and when a saved SBML model is loaded, exchanges that the file leaves open are re-opened after model_adjustments (which switches all exchanges off), exactly as their own all-corrections cell does. The metric is their precision-recall AUC, where labels are the model growth calls, the score is negative fitness and the positive class is no-growth:

$$AUC_{PR} = \int P(r) dr, \quad P(t) = |\{f < -t\} \cap \{model\ no-growth\}| / |\{f < -t\}|, \quad R(t) = |\{f < -t\} \cap \{no-growth\}| / |\{no-growth\}| \quad (H1)$$

Supplements are added as intracellular exchange reactions, the same construction as their Part 6. Our automated set is every cofactor/vitamin-class biomass compound that the iJO1366 provenance audit found unproducible in at least one lethal knockout. It was derived without any fitness data.

variant	PR-AUC	ROC-AUC	no-growth calls	genes x carbons
INVALID: same bug; btn/thmpp/coa re-added the removed vitamins	0.839	0.988	1868	1326 x 25
Bernstein all-corrections model, their saved vitamins kept (published SOTA, reproduced)	0.838	0.988	1868	1326 x 25
INVALID: all-corrections file with its vitamin exchanges closed, loader bug	0.764	0.960	2293	1326 x 25
iML1515 + 5 hand-picked vitamins (Bernstein Part 6)	0.693	0.982	1621	1326 x 25
iJR904 (their pipeline)	0.659	0.921	1846	712 x 25
iAF1260 (their pipeline)	0.638	0.941	1976	1056 x 25
iJO1366 (their pipeline)	0.621	0.920	2221	1148 x 25

variant	PR-AUC	ROC-AUC	no-growth calls	genes x carbons
iML1515 as released (their pipeline)	0.613	0.941	2171	1326 x 25
iML1515 + 19 audit-derived cofactors (ours, label-free)	0.548	0.977	1249	1326 x 25

Table 19. Bernstein benchmark results (their code, their metric).




Figure 11. PR-AUC on the Bernstein benchmark by model variant.

H.3 Reading

Reproduction with the authors' code (PR-AUC, iML1515: 1,326 genes x 25 carbon sources = 33,150 pairs): iML1515 as released 0.613; + the five Bernstein vitamins 0.693; Bernstein all-corrections model with its saved vitamin supplies 0.839 (published 0.843; the 0.004 gap is attributed to solver/library versions and is reported, not hidden). Other releases in the same pipeline: iJO1366 0.621, iAF1260 0.638, iJR904 0.659 (different matched gene sets, so not directly comparable). Our label-free 19-compound cofactor set added to iML1515 scored 0.548: blanket supplementation removes true essentiality calls (no-growth pairs fall from 2,171 to 1,249) and is a NEGATIVE. No SOTA claim is made. The pre-registered split-half search for additional supplements on top of the corrected all-corrections model is reported in the ledger when complete.

H.4 A reproduction error, kept on the record

Our first run of the released all-corrections model scored 0.764. The cause was our loader: it closed every exchange in the saved file, and their model_adjustments also switches exchanges off, so the three vitamin uptakes the authors saved open (EX_btn_e, EX_thm_e, EX_pnto__R_e) were silently removed. Biotin, thiamin and pantothenate synthesis knockouts (b0775, b0776, b3990, b0133, b0134) then showed zero growth on every carbon source. A split-half search over our 19 audit-derived

cofactors selected biotin, thiamin diphosphate and CoA and appeared to beat the model by +0.071 PR-AUC on held-out carbons (95% CI 0.037-0.111). That gain only restored the removed vitamins: after the fix the all-corrections model scores 0.839 on its own, identical to the "selected" run. The rows marked INVALID in Table 19 are kept so the error stays auditable. Lesson: an apparent benchmark win that re-discovers the benchmark authors' own fix is a symptom of a broken baseline.

H.5 Protein language model feature (pre-registered negative)

ESM-2 (t6, 8M parameters; fair-esm) mean-pooled embeddings were computed for all 4,300 CDS of U00096.3. Pre-registered test (notes/prereg_esm_v4.md): v2 features plus a 16-component PCA of the embedding fit inside each training fold. AUROC 0.793 vs 0.797 for v2; paired bootstrap difference -0.004 (95% CI -0.019 to +0.011). The embedding alone reaches AUROC 0.662: real signal, but redundant with the existing features. Verdict: negative.

H.6 Where the ensemble is confidently wrong: GO enrichment

Hard errors of the v2 ensemble were tested for Gene Ontology enrichment (goatools, propagated annotations from the EcoCyc GAF, universe = annotated benchmark genes) with the one-sided hypergeometric test and Benjamini-Hochberg correction:

$$p = \sum_{i \geq k} C(K, i) C(N-K, n-i) / C(N, n), \quad q_{(j)} = \min_{l \geq j} (m p_{(l)} / l) \quad (H2)$$

error set	GO term	study	universe	BH q
confident FP	biosynthetic process	38/45	447/1237	5.6e-08
confident FP	Mo-molybdopterin cofactor biosynthetic process	6/45	11/1237	2.2e-04
confident FP	molybdopterin cofactor metabolic process	6/45	11/1237	2.2e-04
confident FP	Mo-molybdopterin cofactor metabolic process	6/45	11/1237	2.2e-04
confident FP	molybdopterin cofactor biosynthetic process	6/45	11/1237	2.2e-04
confident FP	organophosphate biosynthetic process	19/45	165/1237	2.2e-04
confident FP	ATP biosynthetic process	5/45	7/1237	2.2e-04
confident FP	proton motive force-driven plasma membrane ATP synthesis	5/45	7/1237	2.2e-04

Table 20. GO terms enriched among confident false positives (non-essential genes in the top 10% of v2 scores; n=45). False negatives (n=41): no term at $q < 0.05$.

Molybdopterin cofactor biosynthesis and the ATP synthase are required only anaerobically or are bypassed by fermentation in rich medium; the ensemble inherits a condition mismatch from its FBA inputs, the same class of error Bernstein et al. fixed with vitamins. This is a hypothesis for condition-aware features, not a claim.